Value	Count	Percent
Mr.	517	58.025%
Miss.	185	20.763%
Mrs.	125	14.029%
Master.	40	4.489%
Dr.	7	0.786%
Rev.	6	0.673%
Sir.	5	0.561%
Col.	2	0.224%
Jonkheer.	1	0.112%
Lady.	1	0.112%
the Countess.	1	0.112%
Ms.	1	0.112%

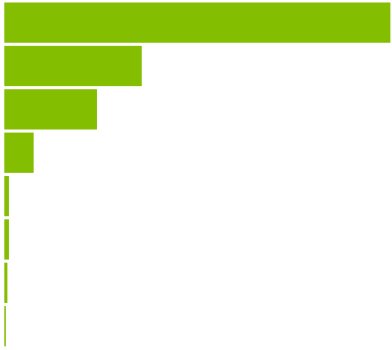

Kaggle
Copyright 2016 Jeff Abrahamson, for private use by course students only

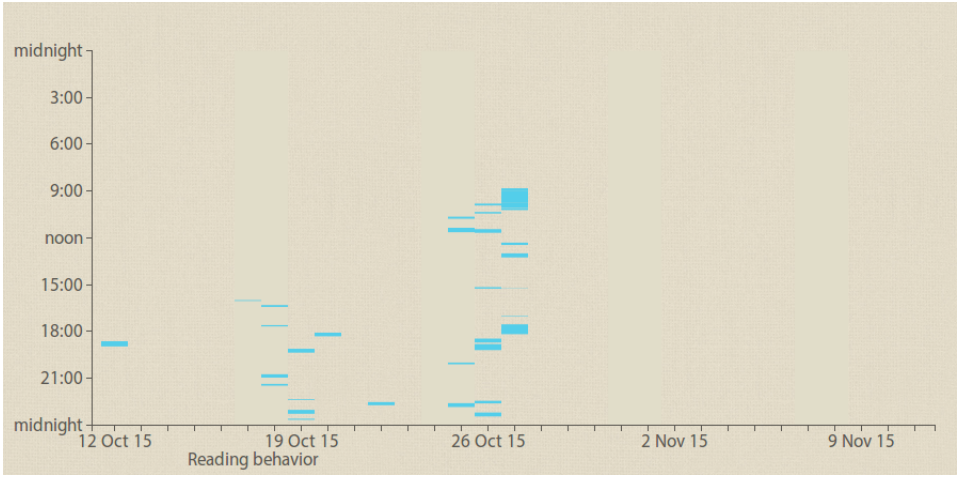

Jellybooks

Copyright 2016 Jeff Abrahamson, for private use by course students only

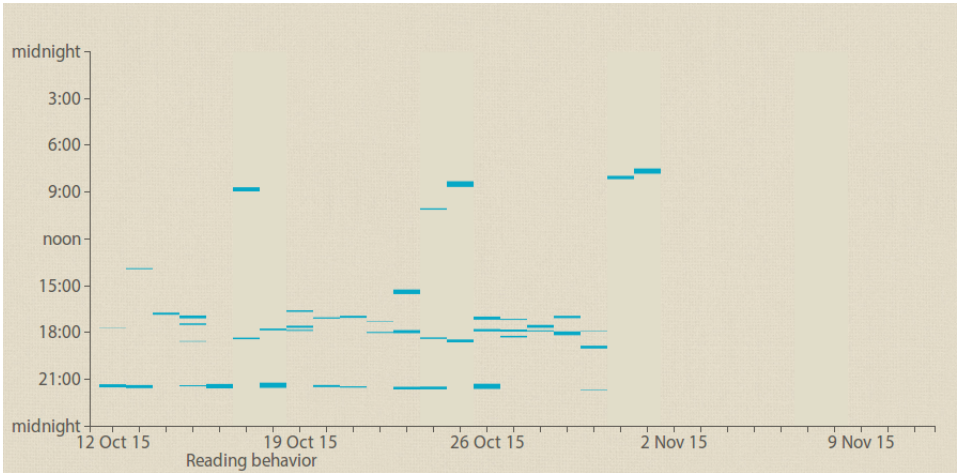

Jellybooks

Copyright 2016 Jeff Abrahamson, for private use by course students only

pandas

```
import pandas as pd
import numpy as np
import scipy
import matplotlib.pyplot as plt
```

Copyright 2016 Jeff Abrahamson, for private use by course students only

pandas

`Dataframe` has many constructors. For example,

```
In [5]: pd.DataFrame({ 'A' : 1.,
                       'B' : pd.Timestamp('20161209'),
                       'C' : pd.Series(1,index=list(range(4)),dtype='float32'),
                       'D' : np.array([3] * 4, dtype='int32'),
                       'E' : pd.Categorical(["test","train","test","train"]),
                       'F' : 'hello' })
Out[5]:
   A          B  C  D      E      F
0  1 2016-12-09  1  3   test  hello
1  1 2016-12-09  1  3  train  hello
2  1 2016-12-09  1  3   test  hello
3  1 2016-12-09  1  3  train  hello

In [6]:
```

Copyright 2016 Jeff Abrahamson, for private use by course students only

Viewing data

```
In [16]: dates = pd.date_range('20161209', periods=4, freq='1w')

In [17]: df = pd.DataFrame(np.random.randn(4,5), index=dates,
                           columns=list('ABCDE'))

In [18]: df.head()
Out[18]:
                   A         B         C         D         E
2016-12-11 -1.303610 -1.235823  0.621914  0.379340 -0.326934
2016-12-18 -1.218197 -1.113826  0.546314 -0.255001 -0.135573
2016-12-25 -0.124625  0.337268 -0.406295  0.587049 -0.904906
2017-01-01 -0.283182 -0.866213  0.051509  0.693037 -0.661055

In [19]:
```

Copyright 2016 Jeff Abrahamson, for private use by course students only

Basic data exploration

```
In [19]: df.describe()
Out[19]:
              A         B         C         D         E
count  4.000000  4.000000  4.000000  4.000000  4.000000
mean  -0.732403 -0.719648  0.203361  0.351106 -0.507117
std    0.614672  0.721194  0.478728  0.424558  0.342755
min   -1.303610 -1.235823 -0.406295 -0.255001 -0.904906
25%   -1.239550 -1.144325 -0.062942  0.220755 -0.722018
50%   -0.750689 -0.990019  0.298912  0.483195 -0.493995
75%   -0.243543 -0.565343  0.565214  0.613546 -0.279094
max   -0.124625  0.337268  0.621914  0.693037 -0.135573

In [20]:
```

Copyright 2016 Jeff Abrahamson, for private use by course students only

Select a column (series)

```
In [20]: df.loc[dates[1]]
Out[20]:
A   -1.218197
B   -1.113826
C    0.546314
D   -0.255001
E   -0.135573
Name: 2016-12-18 00:00:00, dtype: float64

In [21]:
```

Copyright 2016 Jeff Abrahamson, for private use by course students only

Select a range

```
In [21]: df.loc[:, ['A', 'C']]
Out[21]:
                   A         C
2016-12-11 -1.303610  0.621914
2016-12-18 -1.218197  0.546314
2016-12-25 -0.124625 -0.406295
2017-01-01 -0.283182  0.051509

In [22]:
```

Copyright 2016 Jeff Abrahamson, for private use by course students only

Boolean selection criteria

```
In [23]: df[df.D > 0]
Out[23]:
                   A         B         C         D         E
2016-12-11 -1.303610 -1.235823  0.621914  0.379340 -0.326934
2016-12-25 -0.124625  0.337268 -0.406295  0.587049 -0.904906
2017-01-01 -0.283182 -0.866213  0.051509  0.693037 -0.661055

In [24]:
```

Copyright 2016 Jeff Abrahamson, for private use by course students only

Recommended

http://www.gregreda.com/2013/10/26/intro-to-pandas-data-structures/

Copyright 2016 Jeff Abrahamson, for private use by course students only

Plotting

Draw a line

```
import matplotlib.pyplot as plt
plt.plot([1,2,3,4])
plt.ylabel('some numbers')
plt.show()
```

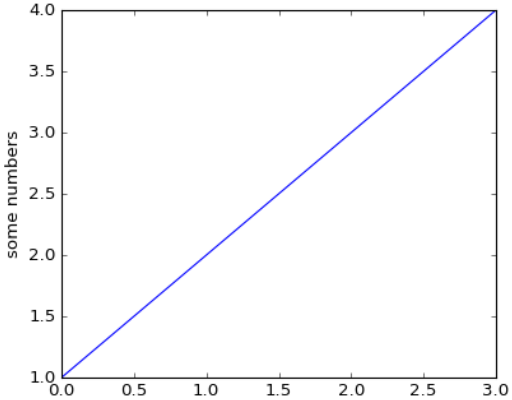

Copyright 2016 Jeff Abrahamson, for private use by course students only

Plotting

Draw a line

```
import matplotlib.pyplot as plt
plt.plot([1, 2, 3, 4], [1, 4, 9, 16])
plt.ylabel('some numbers')
plt.show()
```

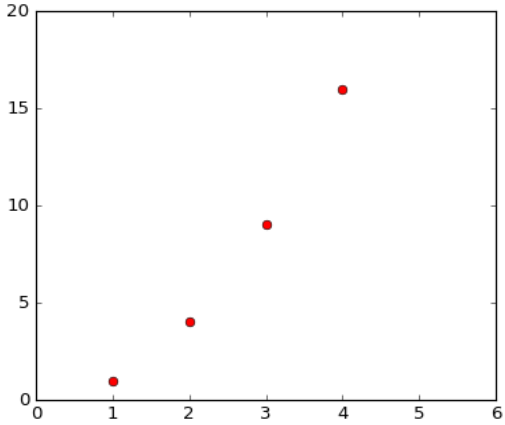

Copyright 2016 Jeff Abrahamson, for private use by course students only

Plotting

Draw a line

```
import numpy as np
import matplotlib.pyplot as plt

# evenly sampled time at 200ms intervals
t = np.arange(0., 5., 0.2)

# red dashes, blue squares and green triangles
plt.plot(t, t,
         'r--', t,
         t**2, 'bs',
         t, t**3, 'g^')
plt.show()
```

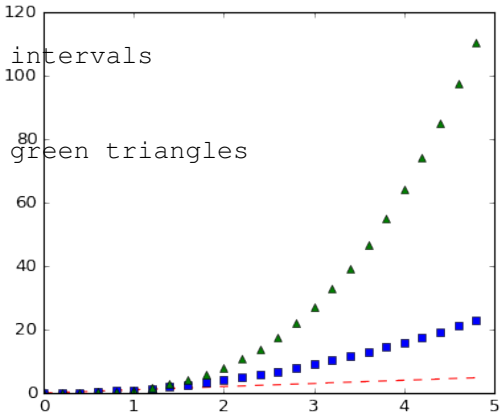

Copyright 2016 Jeff Abrahamson, for private use by course students only

Plotting

Draw two curves

```
import numpy as np
import matplotlib.pyplot as plt

def f(t):
    return np.exp(-t) * np.cos(2*np.pi*t)

t1 = np.arange(0.0, 5.0, 0.1)
t2 = np.arange(0.0, 5.0, 0.02)

plt.figure(1)
plt.subplot(211)
plt.plot(t1, f(t1), 'bo', t2, f(t2), 'k')

plt.subplot(212)
plt.plot(t2, np.cos(2*np.pi*t2), 'r--')
plt.show()
```

Copyright 2016 Jeff Abrahamson, for private use by course students only

Plotting

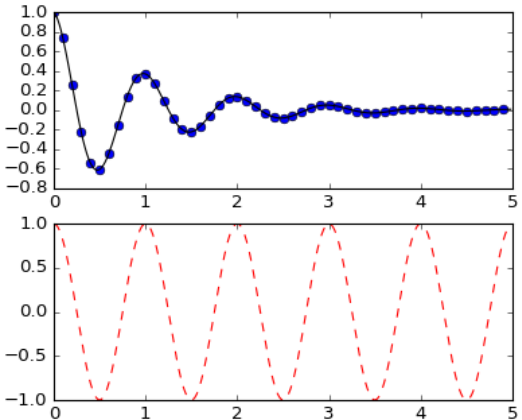

Copyright 2016 Jeff Abrahamson, for private use by course students only

Plotting

Draw two curves

```
import numpy as np
import matplotlib.pyplot as plt

mu, sigma = 100, 15
x = mu + sigma * np.random.randn(10000)

n, bins, patches = plt.hist(x, 50, normed=1, facecolor='g', alpha=0.75)

plt.xlabel('Smarts')
plt.ylabel('Probability')
plt.title('Histogram of IQ')
plt.text(60, .025, r'$\mu=100,\ \sigma=15$')
plt.axis([40, 160, 0, 0.03])
plt.grid(True)
plt.show()
```

Copyright 2016 Jeff Abrahamson, for private use by course students only

Plotting

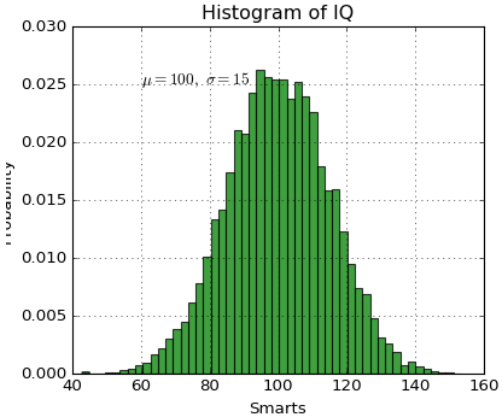

Copyright 2016 Jeff Abrahamson, for private use by course students only

Plotting

Scatter plot

http://matplotlib.org/mpl_examples/pylab_examples/scatter_demo2.py

```
import numpy as np
import matplotlib.pyplot as plt

fig, ax = plt.subplots()
ax.scatter(delta1[:-1], delta1[1:], c=close, s=volume, alpha=0.5)

ax.set_xlabel(r'$\Delta_i$', fontsize=20)
ax.set_ylabel(r'$\Delta_{i+1}$', fontsize=20)
ax.set_title('Volume and percent change')

ax.grid(True)
fig.tight_layout()

plt.show()
```

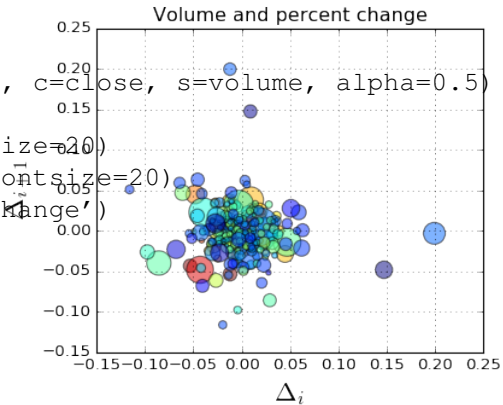

Copyright 2016 Jeff Abrahamson, for private use by course students only

Plotting

http://matplotlib.org/users/pyplot_tutorial.html

http://matplotlib.org/users/beginner.html

Copyright 2016 Jeff Abrahamson, for private use by course students only

Linear models

Problem: $\{(x_i, y_i)\}$.

Given x, predict $\hat{y}$.

Copyright 2016 Jeff Abrahamson, for private use by course students only

Linear models

x: **explanatory** or **predictor** variable.

y: **response** variable.

For some reason, we believe a linear model is a good idea.

Copyright 2016 Jeff Abrahamson, for private use by course students only

Linear models

Example:

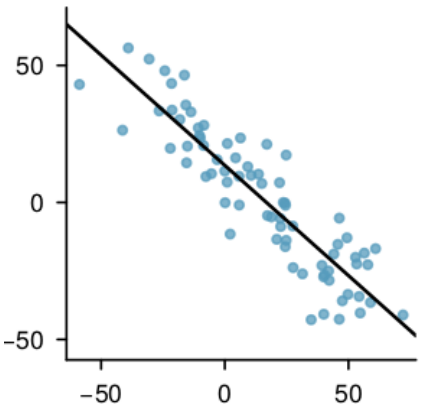

Copyright 2016 Jeff Abrahamson, for private use by course students only

Linear models

Example:

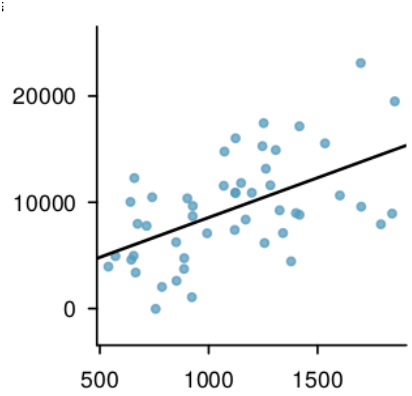

Copyright 2016 Jeff Abrahamson, for private use by course students only

Linear models

Example:

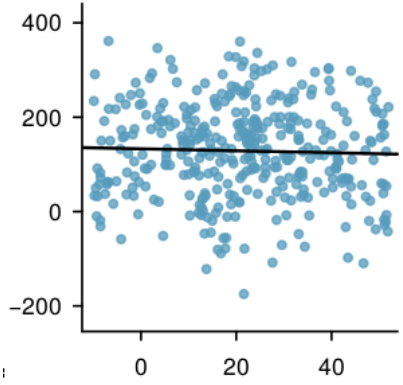

Copyright 2016 Jeff Abrahamson, for private use by course students only

Linear models

Example:

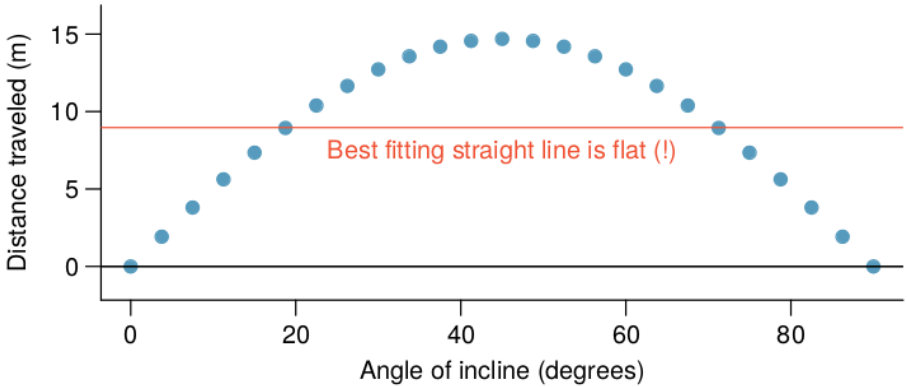

Copyright 2016 Jeff Abrahamson, for private use by course students only

Residuals

What's left over.

$$\text{data} = \text{fit} + \text{residual}$$

Copyright 2016 Jeff Abrahamson, for private use by course students only

Residuals

What's left over.

$$y_i = \hat{y}_i + e_i$$

Copyright 2016 Jeff Abrahamson, for private use by course students only

Residuals

What's left over.

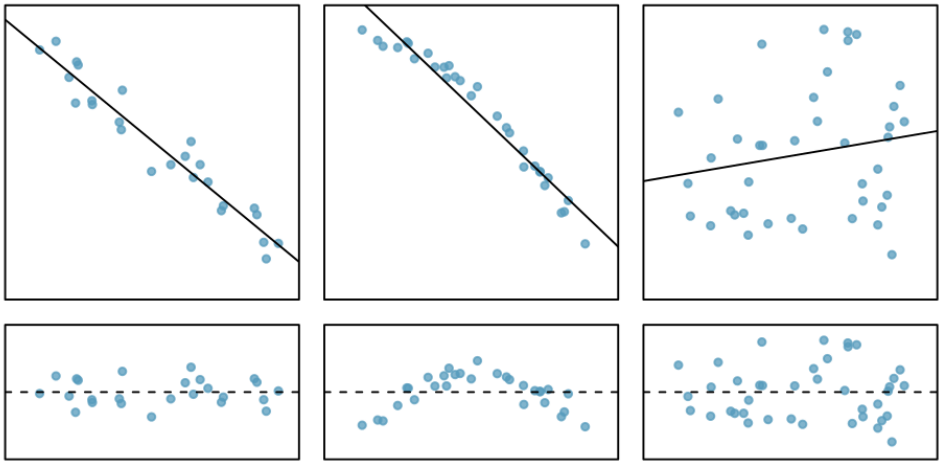

Copyright 2016 Jeff Abrahamson, for private use by course students only

Residuals

What's left over.

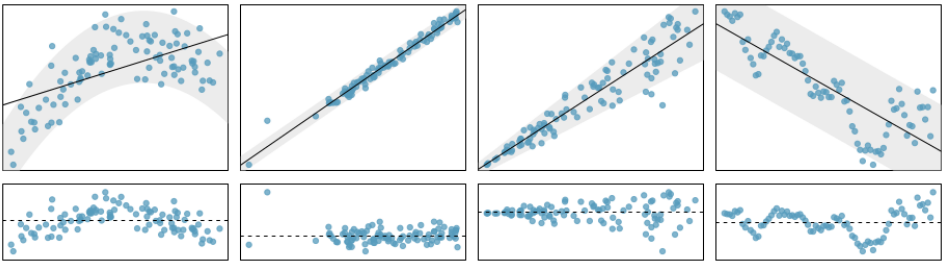

Copyright 2016 Jeff Abrahamson, for private use by course students only

Residuals

What's left over.

Goal: small residuals.

$$\sum | e_i |$$

Copyright 2016 Jeff Abrahamson, for private use by course students only

Residuals

What's left over.

Goal: small residuals.

$$\sum e_i^2$$

Copyright 2016 Jeff Abrahamson, for private use by course students only

Outliers

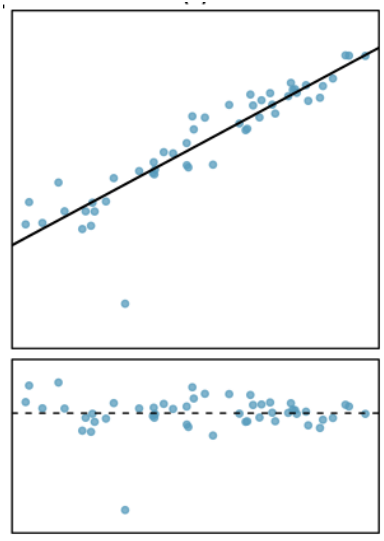

Copyright 2016 Jeff Abrahamson, for private use by course students only

Outliers

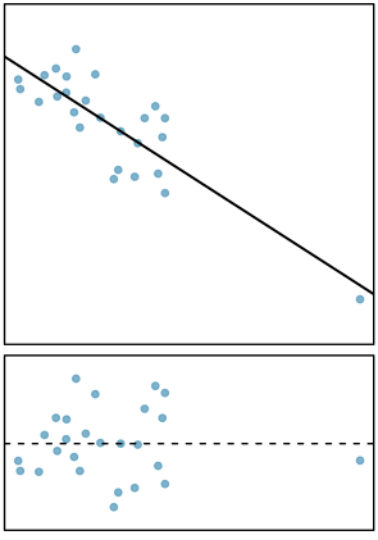

Copyright 2016 Jeff Abrahamson, for private use by course students only

Outliers

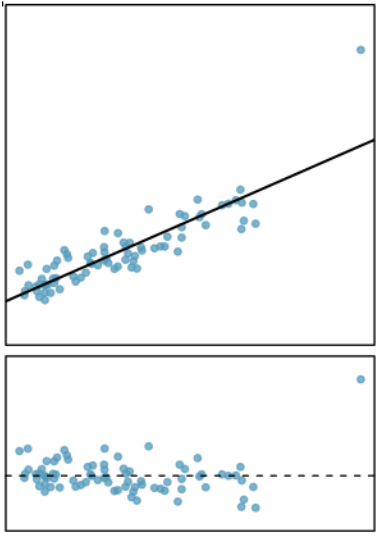

Copyright 2016 Jeff Abrahamson, for private use by course students only

Outliers

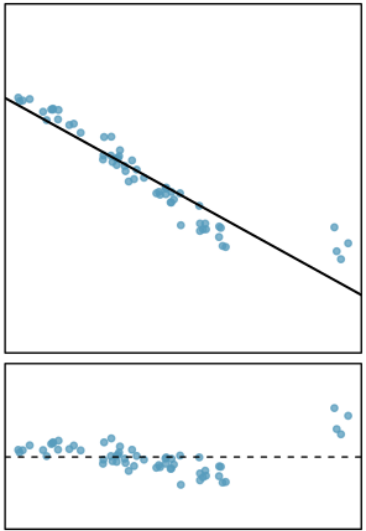

Copyright 2016 Jeff Abrahamson, for private use by course students only

Outliers

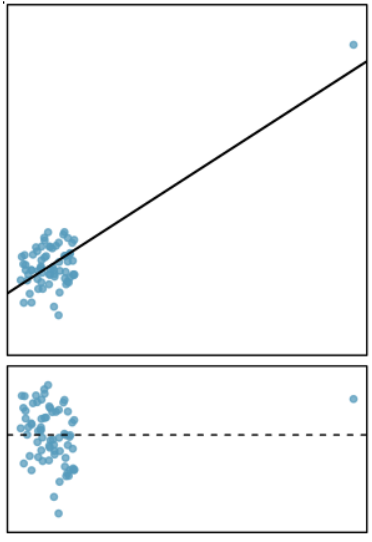

Copyright 2016 Jeff Abrahamson, for private use by course students only

Outliers

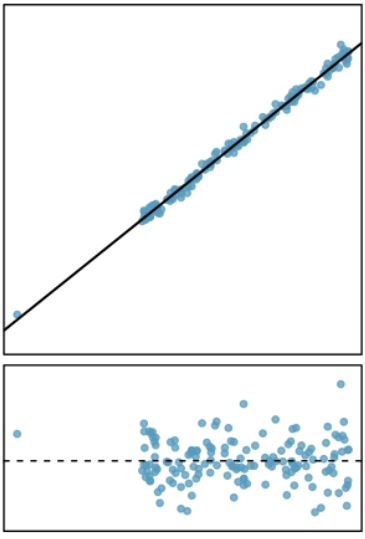

Copyright 2016 Jeff Abrahamson, for private use by course students only

Outliers

Don’t ignore outliers.

Copyright 2016 Jeff Abrahamson, for private use by course students only

Outliers

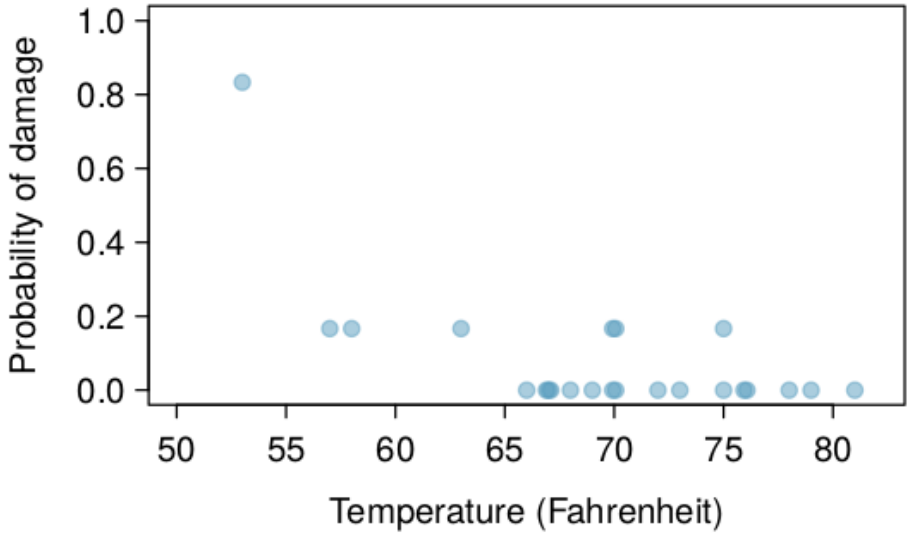

Copyright 2016 Jeff Abrahamson, for private use by course students only

Correlation

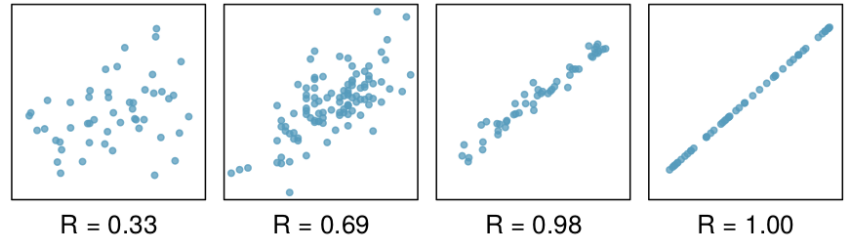

Copyright 2016 Jeff Abrahamson, for private use by course students only

Correlation

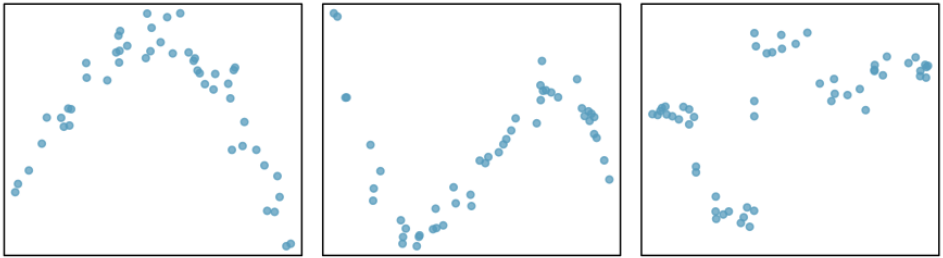

Copyright 2016 Jeff Abrahamson, for private use by course students only

Correlation

Anscombe's Quartet

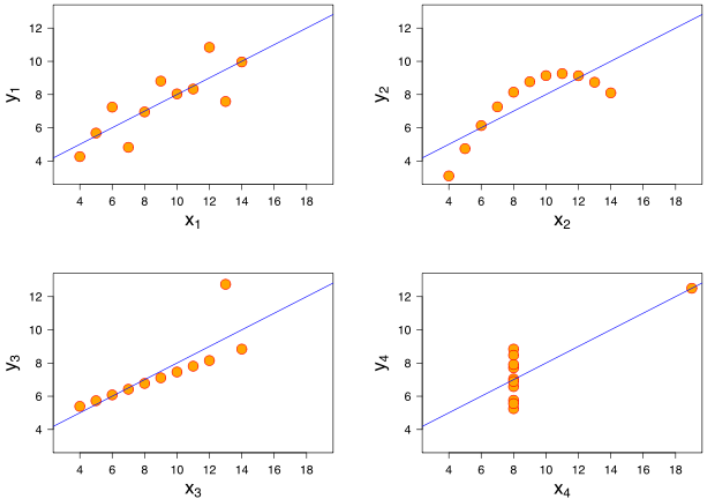

Copyright 2016 Jeff Abrahamson, for private use by course students only

Correlation does not imply causation

Copyright 2016 Jeff Abrahamson, for private use by course students only

Hypothesis (model)

$$h_\theta(x) = \theta_0 + \theta_1 x$$

Copyright 2016 Jeff Abrahamson, for private use by course students only

Cost function

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^{m} (h_\theta(x_i) - y_i)^2$$

Copyright 2016 Jeff Abrahamson, for private use by course students only

Gradient descent

$$\begin{cases} \theta_0 & \leftarrow \theta_0 - \alpha \dfrac{\partial}{\partial \theta_0} J(\theta_0, \theta_1) \\ \theta_1 & \leftarrow \theta_1 - \alpha \dfrac{\partial}{\partial \theta_1} J(\theta_0, \theta_1) \end{cases}$$

Copyright 2016 Jeff Abrahamson, for private use by course students only

Gradient descent

$$\begin{cases} \theta_0 & \leftarrow \theta_0 - \dfrac{\alpha}{m} \displaystyle\sum_{i=1}^{m} (h_\theta(x_i) - y_i) \\ \theta_1 & \leftarrow \theta_1 - \dfrac{\alpha}{m} \displaystyle\sum_{i=1}^{m} (h_\theta(x_i) - y_i) \end{cases}$$

Copyright 2016 Jeff Abrahamson, for private use by course students only

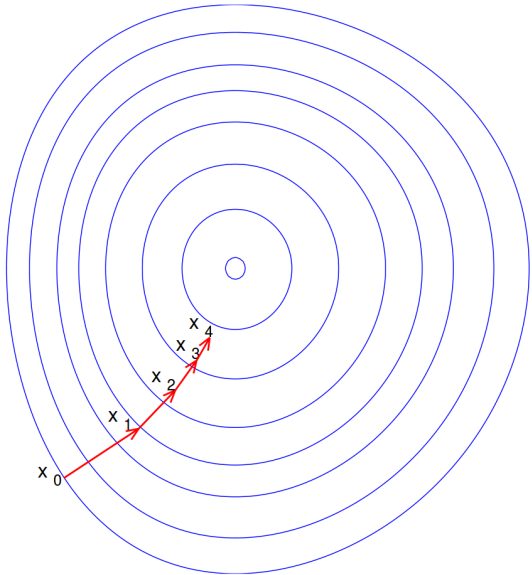

Copyright 2016 Jeff Abrahamson, for private use by course students only

Hypothesis again

$$
\begin{aligned}
h_\theta(x) &= \theta_0 + \theta_1 x_1 \\
&= \theta_0 + \sum_{i=1}^{1} \theta_i x_i \\
&= [\theta_0, \theta_1] \begin{bmatrix} x_0 \\ x_1 \end{bmatrix} \\
&= \theta^T x
\end{aligned}
$$

Copyright 2016 Jeff Abrahamson, for private use by course students only

Hypothesis (multiple regression)

$$h_\theta(x) = \theta_0 + \sum_{i=1}^{n} \theta_i x_i$$

$$= [\theta_0, \cdots, \theta_n] \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_n \end{bmatrix}$$

$$= \theta^T x$$

Copyright 2016 Jeff Abrahamson, for private use by course students only

Hypothesis (multiple regression)

$$h_\theta(x) = \theta^T x$$
$$= \theta^T x^{(1)}$$

Copyright 2016 Jeff Abrahamson, for private use by course students only

Hypothesis (multiple regression)

$$X = \begin{bmatrix} | & | & \cdots & | \\ x^{(1)} & x^{(2)} & \cdots & x^{(m)} \\ | & | & \cdots & | \end{bmatrix} = \begin{bmatrix} x_0^{(1)} & x_0^{(2)} & \cdots & x_0^{(m)} \\ x_1^{(1)} & x_1^{(2)} & \cdots & x_1^{(m)} \\ \vdots & \vdots & \ddots & \vdots \\ x_n^{(1)} & x_n^{(2)} & \cdots & x_n^{(m)} \end{bmatrix}$$

Copyright 2016 Jeff Abrahamson, for private use by course students only

Hypothesis (multiple regression)

$$
\begin{aligned}
h_\theta(X) &= \theta^T X \\
&= \left[h_0(x^{(1)}), h_0(x^{(2)}), \cdots, h_0(x^{(m)})\right] \\
&= \theta^T X
\end{aligned}
$$

Copyright 2016 Jeff Abrahamson, for private use by course students only

Hypothesis (multiple regression)

or $X\theta$ if row vectors...

Copyright 2016 Jeff Abrahamson, for private use by course students only

Cost function (multiple regression)

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^{m} \left(h_\theta(x^{(i)}) - y^{(i)} \right)^2$$
$$= \frac{1}{2m} (X\theta - Y)^T (X\theta - Y)$$

Copyright 2016 Jeff Abrahamson, for private use by course students only

Gradient descent (multiple regression)

$$\theta_j \leftarrow \theta_j - \frac{\alpha}{m} \sum_{i=1}^{m} \left(h_\theta(x^{(i)}) - y^{(i)} \right) \cdot x_j^{(i)}$$

for $j = 1, \cdots, n$

Copyright 2016 Jeff Abrahamson, for private use by course students only

Gradient descent (multiple regression)

$$\theta \leftarrow \theta - \nabla \boldsymbol{J}(\theta)$$

$$\text{where } \nabla = \begin{bmatrix} \frac{\partial}{\partial \theta_0} \\ \frac{\partial}{\partial \theta_1} \\ \vdots \\ \frac{\partial}{\partial \theta_n} \end{bmatrix}$$

Copyright 2016 Jeff Abrahamson, for private use by course students only

Scikit Learn

```
>>> import numpy as np
>>> from sklearn.linear_model import LinearRegression
>>> X = np.array([[1, 1], [1, 2], [2, 2], [2, 3]])
>>> # y = 1 * x_0 + 2 * x_1 + 3
>>> y = np.dot(X, np.array([1, 2])) + 3
>>> reg = LinearRegression().fit(X, y)
>>> reg.score(X, y)
1.0
>>> reg.coef_
array([1., 2.])
>>> reg.intercept_
3.0000...
>>> reg.predict(np.array([[3, 5]]))
array([16.])
```

Copyright 2016 Jeff Abrahamson, for private use by course students only