

# Review of Linear Regression

# Linear models

**Problem:**  $\{(x_i, y_i)\}$ .

Given  $x$ , predict  $\hat{y}$ .

# Linear models

$x$ : **explanatory** or **predictor** variable. Or the **signal**.

$y$ : **response** variable.

For some reason, we believe a linear model is a good idea.

# Residuals

Residuals (model error):

- What's left over
- What the model doesn't explain

$$\text{data} = \text{fit} + \text{residual}$$

# Residuals

Residuals (model error):

- What's left over
- What the model doesn't explain

$$y_i = \hat{y}_i + e_i$$

# Residuals

Residuals (model error):

- What's left over
- What the model doesn't explain

$$e_i = y_i - \hat{y}_i$$

# Residuals

Residuals (model error):

- What's left over
- What the model doesn't explain

Goal: small residuals.

$$\sum e_i^2$$

# Hypothesis (Model)

$$y = a + bx$$



# Hypothesis (Model)

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

# Cost Function

Also called the “loss function”.

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (e_i)^2$$

# Cost Function


Also called the “loss function”.

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (\hat{y}_i - y_i)^2$$

# Cost Function

Also called the “loss function”.

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x_i) - y_i)^2$$



# Gradient Descent

$$\begin{cases} \theta_0 & \leftarrow \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1) \\ \theta_1 & \leftarrow \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1) \end{cases}$$

# Gradient Descent

$$\begin{cases} \theta_0 \leftarrow \theta_0 - \alpha \frac{\partial}{\partial \theta_0} \left( \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x_i) - y_i)^2 \right) \\ \theta_1 \leftarrow \theta_1 - \alpha \frac{\partial}{\partial \theta_1} \left( \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x_i) - y_i)^2 \right) \end{cases}$$

# Gradient Descent

$$\begin{cases} \theta_0 \leftarrow \theta_0 - \alpha \frac{\partial}{\partial \theta_0} \left( \frac{1}{2m} \sum_{i=1}^m (\theta_0 + \theta_1 x_i - y_i)^2 \right) \\ \theta_1 \leftarrow \theta_1 - \alpha \frac{\partial}{\partial \theta_1} \left( \frac{1}{2m} \sum_{i=1}^m (\theta_0 + \theta_1 x_i - y_i)^2 \right) \end{cases}$$



# Gradient Descent

$$\begin{cases} \theta_0 \leftarrow \theta_0 - \frac{\alpha}{m} \sum_{i=1}^m (\theta_0 + \theta_1 x_i - y_i) \cdot 1 \\ \theta_1 \leftarrow \theta_1 - \frac{\alpha}{m} \sum_{i=1}^m (\theta_0 + \theta_1 x_i - y_i) x_i \end{cases}$$

# Gradient Descent

$$\begin{cases} \theta_0 \leftarrow \theta_0 - \frac{\alpha}{m} \sum_{i=1}^m (h_{\theta}(x_i) - y_i) \\ \theta_1 \leftarrow \theta_1 - \frac{\alpha}{m} \sum_{i=1}^m (h_{\theta}(x_i) - y_i) x_i \end{cases}$$

# Logistic regression

# Linear regression

- Inputs are continuous or discrete
- Continuous output
- Normal residues
- Predict  $\hat{y}$  for  $x$  given  $\{(x_i, y_i)\}$

# Logistic regression

- Inputs are continuous or discrete
- Binary output
- Classification

# Logistic regression

- Have: continuous and discrete inputs
- Want: class (0 or 1)

# Logistic regression: motivation

# Probabilistic inspiration

The probabilities are motivations: this doesn't really behave like a probability.

$$h_{\theta}(x) = .75 \iff \text{event has 75\% of being true}$$



# Probabilistic inspiration

The probabilities are motivations: this doesn't really behave like a probability.

$$h_{\theta}(x) = \Pr(y = 1 \mid x; \theta) = 0.75$$

# Probabilistic inspiration

The probabilities are motivations: this doesn't really behave like a probability.

So this must be true:

$$\Pr(y = 0 \mid x; \theta) + \Pr(y = 1 \mid x; \theta) = 1$$

# Probabilistic inspiration

The probabilities are motivations: this doesn't really behave like a probability.

$$\text{Set } y = 1 \iff h_{\theta}(x) = \Pr(y = 1 \mid x; \theta) > \frac{1}{2}$$

# Probabilistic inspiration

The probabilities are motivations: this doesn't really behave like a probability.

Math review:

- $z = (\theta^T x)$
- $\theta^T x \geq 0 \iff h_\theta \geq 0.5$
- $\theta^T x \geq 0 \iff \text{predict } y = 1$

# Logistic (sigmoid, logit) function

$$\sigma(z) = \frac{e^z}{e^z + 1} = \frac{1}{1 + e^{-z}}$$

# Logistic (sigmoid, logit) function

$$\sigma(z) = \frac{e^z}{e^z + 1} = \frac{1}{1 + e^{-z}}$$

Exercise: plot this

# Logistic (sigmoid, logit) function

$$\sigma(z) = \frac{e^z}{e^z + 1} = \frac{1}{1 + e^{-z}}$$

Let

$$z = h_{\theta}(x) = \theta_0 + \theta_1 x$$

# Logistic (sigmoid, logit) function

$$\sigma(z) = \frac{e^z}{e^z + 1} = \frac{1}{1 + e^{-z}}$$

Then

$$\sigma(z) = \frac{1}{1 + e^{-(\theta_0 + \theta_1 x)}}$$



# Cost function in logistic regression

In linear regression, we had

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (e_i)^2$$

# Cost function in logistic regression

In linear regression, we had

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (\hat{y}_i - y_i)^2$$

# Cost function in logistic regression

In linear regression, we had

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x_i) - y_i)^2$$

# Cost function in logistic regression

Here's a convex cost function:

$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$

# Cost function in logistic regression

Here's a convex cost function:

$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$

Exercise: Plot this (cost vs  $y$ ).

# Cost function in logistic regression

Here's a convex cost function:

$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x), y)$$

# Cost function in logistic regression

Here's a convex cost function:

$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$

$$J(\theta) = y \cdot \log(h_{\theta}(x)) + (1 - y) \cdot \log(1 - h_{\theta}(x))$$

# Gradient descent

$$\theta_j \leftarrow \theta_j - \frac{\alpha}{m} \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right) \cdot x_j^{(i)}$$

for  $j = 1, \dots, n$



# null hypothesis

**true positive, true negative**

**false positive, false negative**

## **type I error**

(incorrect rejection of null hypothesis)

## **type II error**

(failure to reject null hypothesis)

## **sensitivity**

100% sensitivity = no false negatives

## **specificity**

100% specificity = no false positives

# Precision

$$P = \frac{TP}{TP + FP}$$

# Recall

$$R = \frac{TP}{TP + FN}$$

# F1 score

$$F1 = \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$



# Non-linear decision boundaries

non-linear-boundary.png

# Non-linear decision boundaries

$$\text{OvA} = \text{OvR}$$

$$\text{OvO}$$

# Non-linear decision boundaries

One vs All = One vs Rest

One vs One